

Week 2 Workshop: Setting up reproducible data science with R

Charlotte Hadley

1. Reproducible code

2. Reproducible referencing

3. Working with datasets in R reproducibly

4. Understanding the assessment template

Reproducible code

Research code is research software

If your research includes R code that reads and wrangles your data in such a way that you can:

- Create charts
- Test statistical hypotheses & compute statistical measures
- Calculate measures of your data

... then you're writing research software and it's an intrinsic part of your research.

You might even develop your R code into an R package that others can reuse for similar analyses. However, that's definitely beyond the scope of this course.

Now you're writing software we need to equip you with some software development tools

Version control with git

Version control is extremely important in software development to ensure future changes to code aren't breaking code.

This is particularly true when code is being worked on by multiple people.

Version control with git

Version control is extremely important in software development to ensure future changes to code aren't break existing code.


This is particularly true when code is being worked on by multiple people.

The most widely used version control tool used today is called `git`.

`git` is an extremely powerful, open-source tool.

We are going to ignore almost all of its features and use only what we need.

Just like most people who use it.



<https://xkcd.com/1597/>

git repositories

Example project that uses git is called a repository.

If you'd like to use git to version control your assignment then that would be a repository.

git and the command line

git was originally designed to be used in the command line

- On Windows that usually means using cmd.exe
- On macOS and Linux that means the Terminal

This practically means that you'd be using a command line interface (CLI) to work with the git technology.

In the case of git it can be difficult to use a GUI interface instead of the CLI to conceptualise what actions you're about to - and often more importantly - what actions were taken in the past.

GitHub

GitHub is one of many different tools for managing git repositories.

We're going to prefer GitHub because it has lots of really awesome features.

It's also owned by Microsoft which radically simplifies installing git onto your machine.

Let's first create an account.

1. Visit github.com and click sign-up (if you don't already have an account)
2. Don't worry about being creative with your username now - you can change it later without things exploding too much.



GitHub Desktop

When you install GitHub Desktop it installs both `git` and a GUI for working with repositories on GitHub.

1. Download the application from desktop.github.com/
2. Run through the installer
3. Sign-in with your GitHub account.



Organising your computer

From this point onward we need to be careful about organising our code.

- RStudio projects should not contain projects themselves
- Git repositories should never contain another git repository¹


It's best practice to have one folder on your machine call GitHub and for all repositories to be stored in that folder.

[1] This is getting into quite advanced territory, but if you needed a repo to depend on another you would use [submodules](#).

Creating a new repository

We're going to make use of the `{usethis}` package to help us create a well-formed new repository.


Please open RStudio and install the `{usethis}` package.



Creating a new repository

We're going to make use of the `{usethis}` package to help us create a well-formed new repository.

I'm going to demonstrate how to do this first afterwards I'll show the instructions and then create your own repository.



Creating a new repository

1. In RStudio choose the "Project Menu" choose

- New Project...
- New directory
- New project

2. Set the name of the directory as `eng7218-week-2-workshop`

3. Create the project as a subdirectory of your GitHub folder

4. Inside of this project run this in the R console

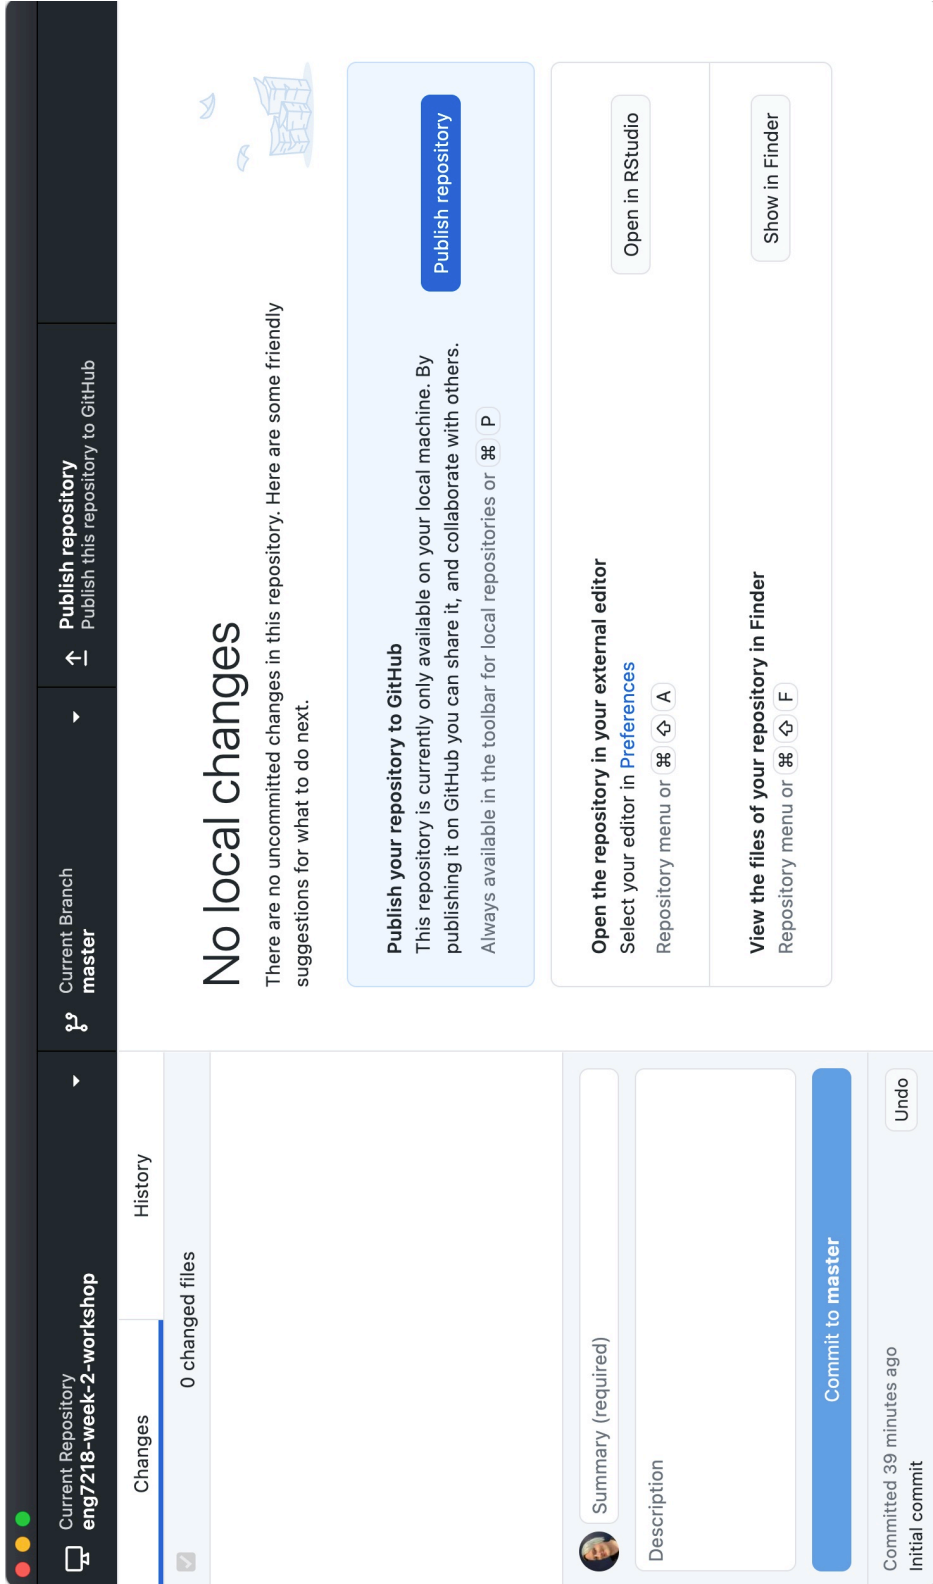
```
use_this::use_git()
```

5. Open GitHub Desktop and select

File > Add Local Repository

Publish to GitHub.com

If everything is good your GitHub Desktop will look like this - let's publish the repository!



Publish to GitHub.com

- We then get to name the repository on GitHub.com - use the same name!!
- We then get to choose if our code is private or not
 - If you're applying an embargo to your data, you probably want to do the same on GitHub!
 - Public repositories can be seen by anyone and are useful for your CV

Let's add some sensitive data to the repo (I)


When adding data to your repository it's important to have two directories .

- data - contains your sensitive data and data anonymisation scripts
- data contains a copy of your anonymous data

We're going to prevent the data - folder from being set to GitHub.com

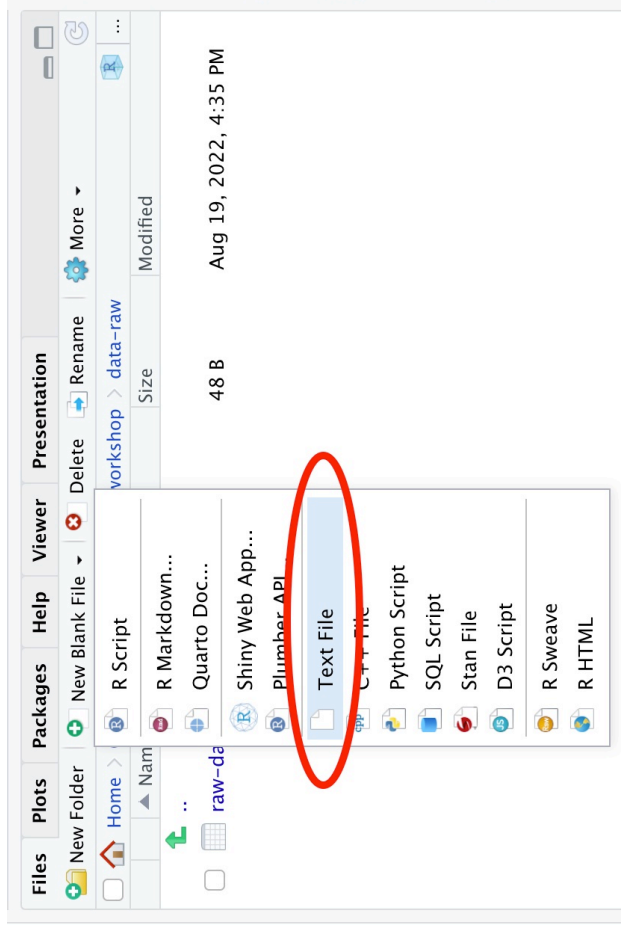
Let's add some sensitive data to the repo (II)

1. Use the New Folder button to create a folder called data - raw



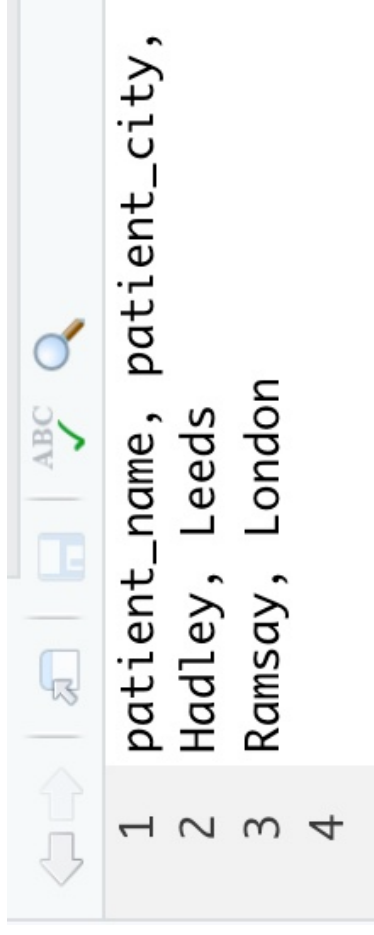
Let's add some sensitive data to the repo (II)

1. Use the New Folder button to create a folder called data - raw in the Files tab
2. Click inside the data - raw folder in the Files tab
3. Select "Text file" from the "New Blank File" menu



Let's add some sensitive data to the repo (II)

1. Use the New Folder button to create a folder called data - raw
2. Click inside the data - raw in the Files tab
3. Select "Text file" from the "New Blank File" menu
4. Add some fake data to the file
5. Save this file into the data - raw and call it something like patient - data . csv



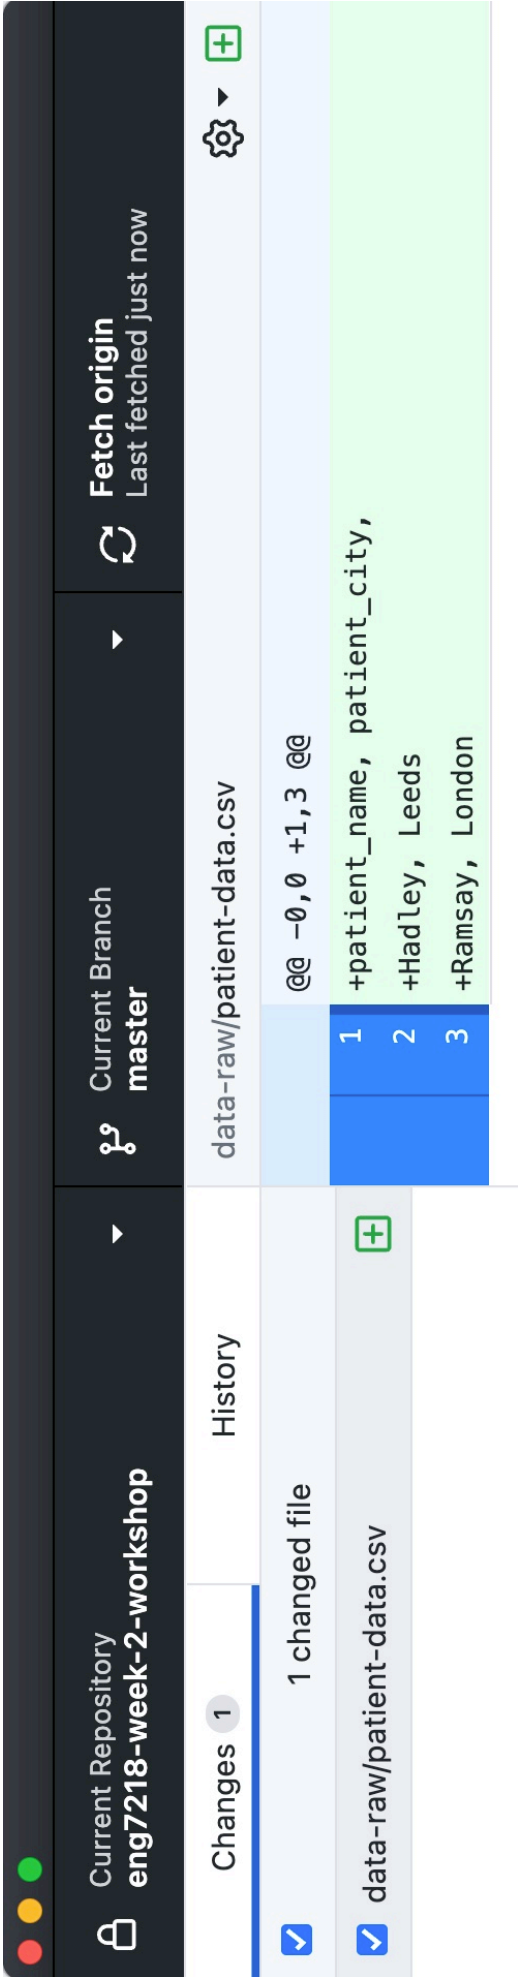
1	patient_name, patient_city,
2	Hadley, Leeds
3	Ramsay, London
4	

Check GitHub Desktop

When you switch to GitHub Desktop we see that git has tracked changes in the repo.

We want to prevent this entire directory from going to GitHub

- The repo on your machine is called "the local repository"
- The repo on Github is called "the remote repository"



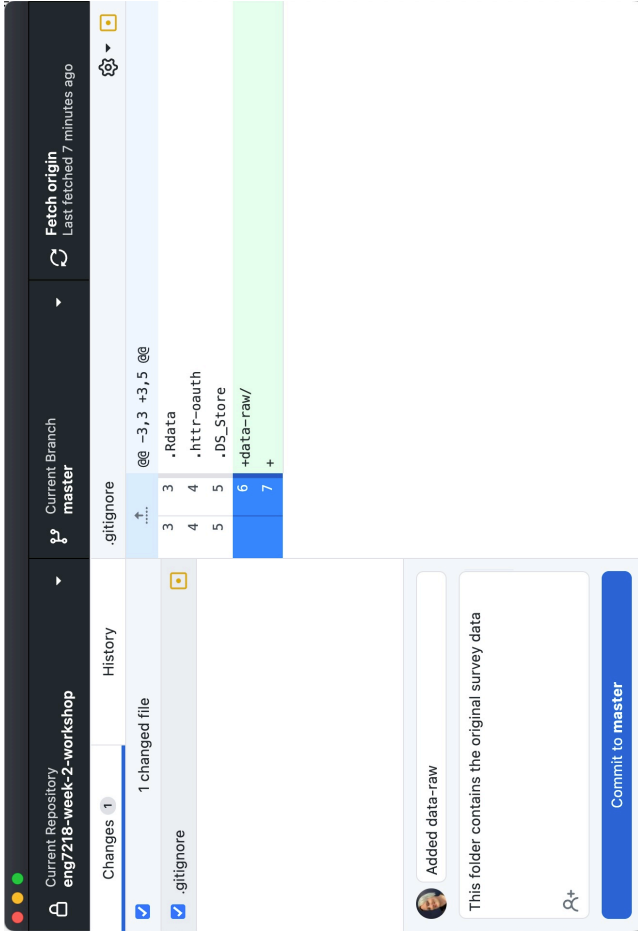
Commit this change to the repository

1. Switch back to GitHub
2. Add the commit messages

These build a branch and sweet, like "added a raw"

3. Optionally add a summary
4. Click "Commit to master"

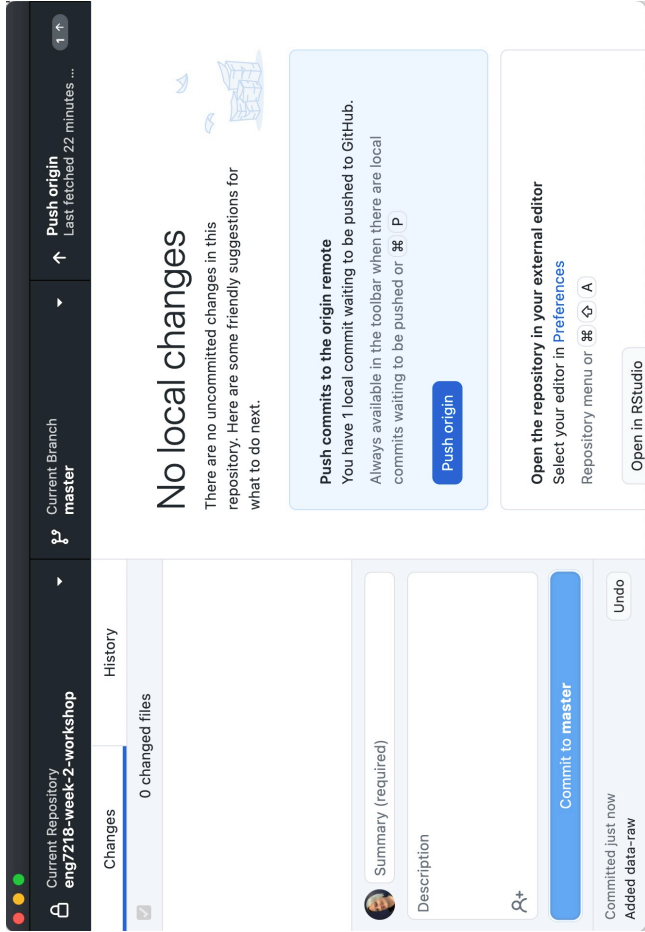
In general, when we "commit" something we tell git to remember our change in the local repository.



Pushing our changes to the remote repository

After committing locally it's time to push the changes to GitHub.com

Click the "push origin" button and the changes will be sent to the remote!

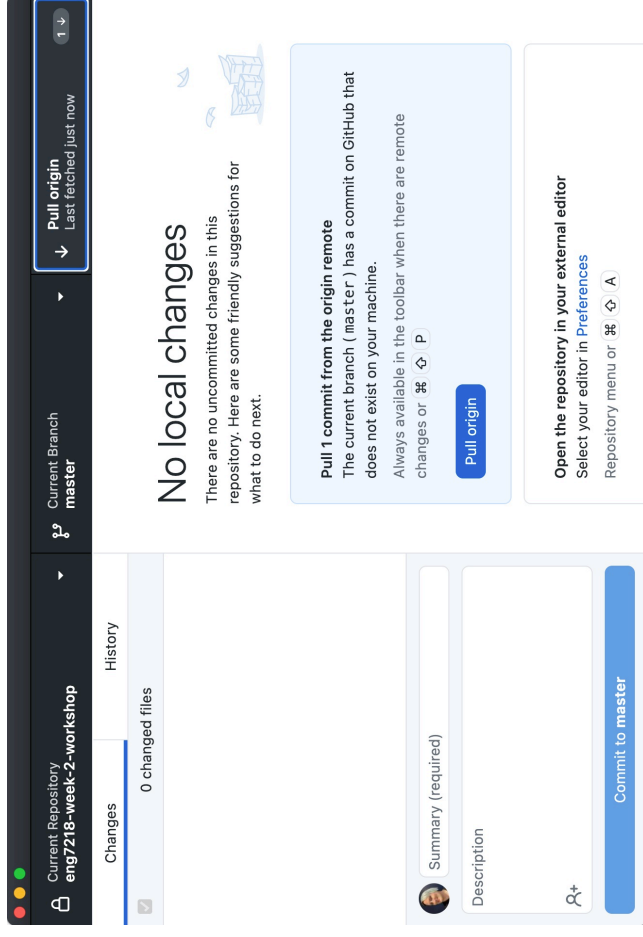


Pulling is the opposite of pushing

Remember that git is designed for collaboration.

Imagine a collaborator has pushed a change to the remote repository but you don't have those changes in your machine.

If you refresh GitHub Desktop you'll be prompted to "pull origin" - which will bring the changes to your local repository.



GitHub branches and pull requests

This is as far as we'll go with using `git` for version control.

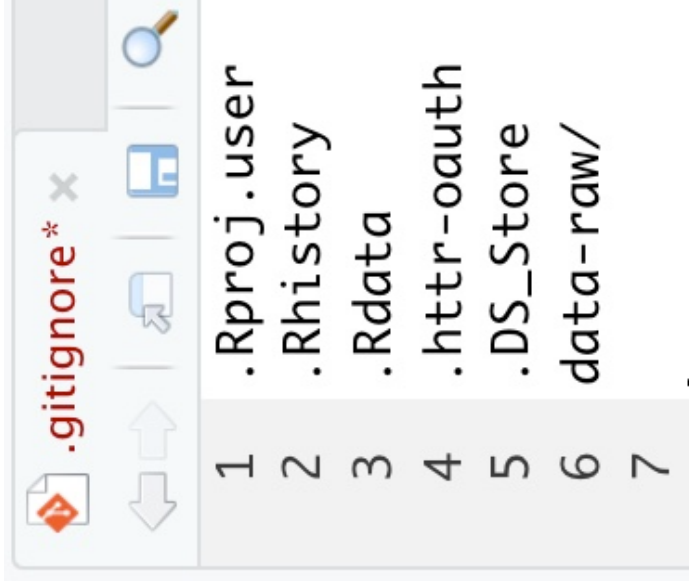
In real-world projects you would use two additional features: branches and pull requests.

There is a really good example to follow tutorial here: <https://datastudies.github.io/hello-world>

There are also tutorials on LinkedIn Learning.

Modify the .gitignore file

1. Switch back to RStudio
2. Go to the project directory
3. Open the .gitignore file
4. Add data - raw to the file and save it



Connecting Figshare & Github

GitHub is an awesome place for code to live.

It's a bad choice for preserving code for the future.

Fortunately, it's simple to connect your Figshare account with GitHub which gives you a fairly reproducible coding workflow.

Reproducible referencing

Referencing is part of research (and the assessment)

Managing your references is an unavoidable part of doing research and is part of the assessment.

There are several reference management tools available.

- EndNote is widely taught however it's close source, expensive and [in my own personal opinion] rubbish
- Mendeley is much better but it's closed source has limited shared library functionality.
- Zotero is open source, has good shared library functionality and has lots of extensions.



I'm going to teach you how to use Zotero.
But choose which software works for you!

Workshop 2 of 8

Setting up Zotero

1. Create a free account at [zotero.org/](https://www.zotero.org/)
2. Download the software from [zotero.org/](https://www.zotero.org/)
3. Install the software and login to your Zotero account

Improving Zotero with Better Bibtex for Zotero

We're going to install the "Better Bibtex" addin which improves how Zotero generates bibtex files.

LaTeX is an extremely useful and very widely used tool for storing, managing and typesetting references

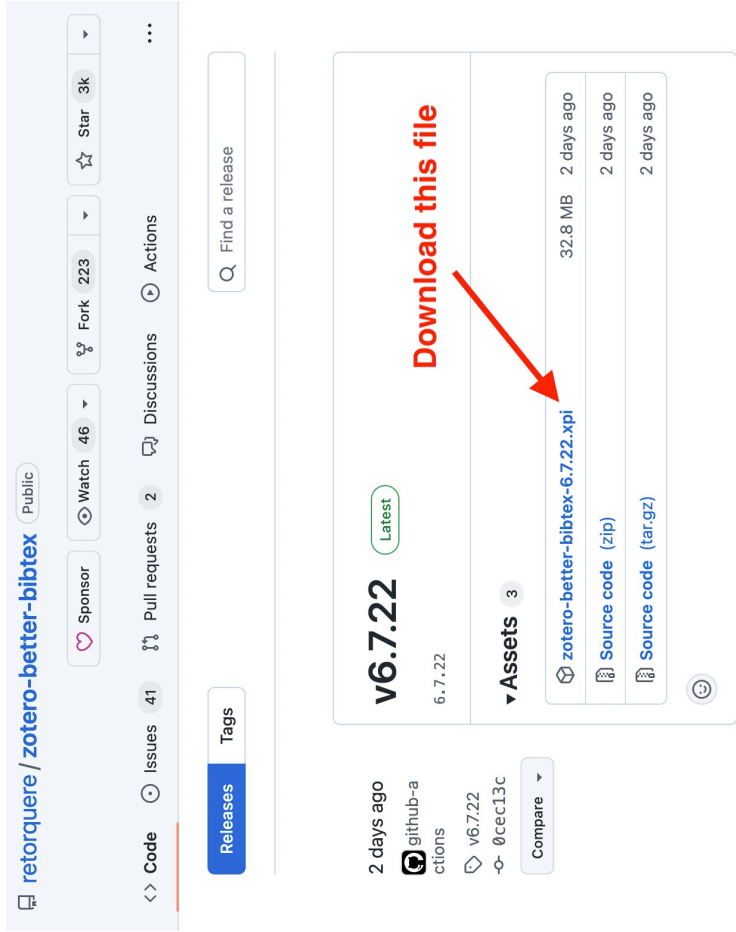
It was originally designed to work with LaTeX which we're purposefully avoiding.

Improving Zotero with Better Bibtex for Zotero

The instructions for installing Better Bibtex.

1. Download the tool from

<https://github.com/retorquere/zotero-better-bibtex/releases>



Improving Zotero with Better Bibtex for Zotero

The instructions for installing Better Bibtex.

1. Download the tool from <https://github.com/retorquere/zotero-better-bibtex/releases>

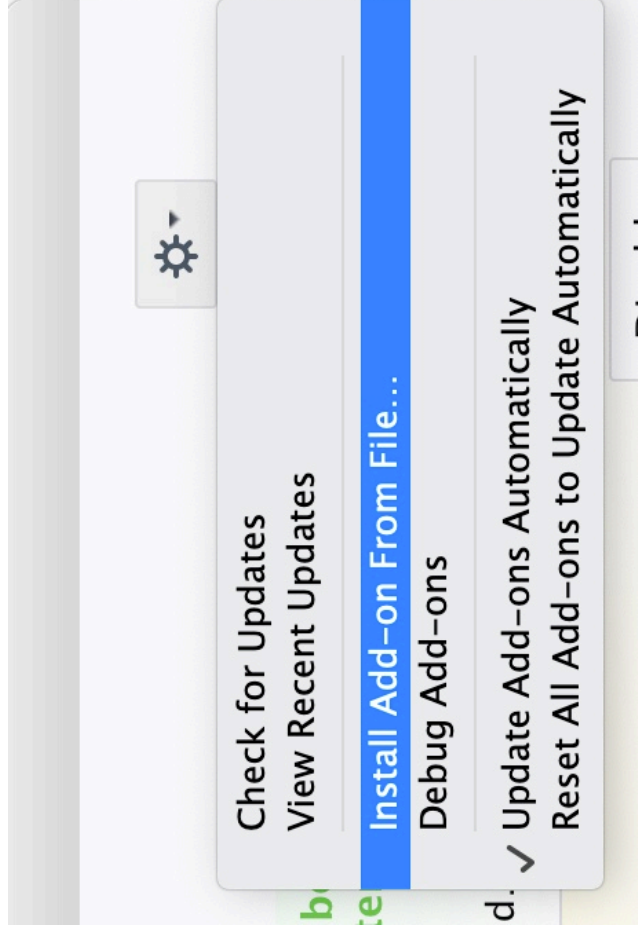
2. Open Zotero and select

Tools > Add-ons

3. In the popup click on the gear icon and select "Install Add-on From File..."

4. Select the .xpi file you just downloaded

5. Restart Zotero





Web browser plug-in

One of the best things about Zotero is how easy it is to add references to your collection directly from the web browser

Please install the plugin for your browser of choice

[zotero.org/download-connectors](https://www.zotero.org/download-connectors)

[Home](#) > Connectors



Adding a paper to Zotero

I'd like you to add the "Why Most Published Research Findings are False" paper to your Zotero collection.

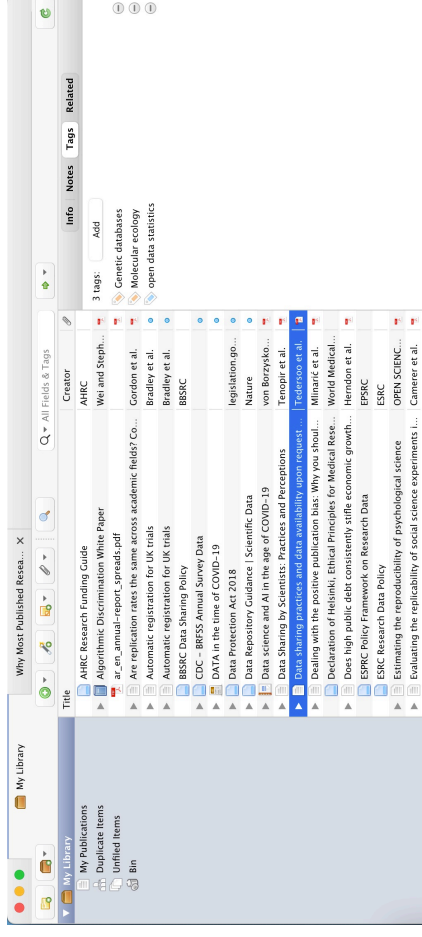
You need your browser and Zotero open simultaneously.

1. Visit <https://doi.org/10.1371/journal.pmed.0020124>
2. Click on the Zotero plugin icon and click "Done"
3. Go over to Zotero and you'll see the paper has been added!

Using Zotero

You can now use Zotero to organise your papers, web pages and reports. All of the things you need to cite and reference.

- Zotero includes a PDF viewer where you can annotate and highlight PDFs.
- You can organise things via folders.
- Tags are another useful tool for organising research



Connecting this to an RStudio project

Our goal is to set up a workflow where we can synchronise our Zotero library with an RStudio project so that we can write a report and add references as we grow our library.

We installed Better BibTeX because it allows us to automate this process.

You'll need to do this for each project you want to connect to your Zotero collection.

Sync Zotero with an RStudio Project

1. Add an R Markdown document to your RStudio

project called something like `Report.Rmd`

- The YAML header is looking for a BibTeX file called `references.bib` which we've not created yet
- We've added a "References" header before the references.

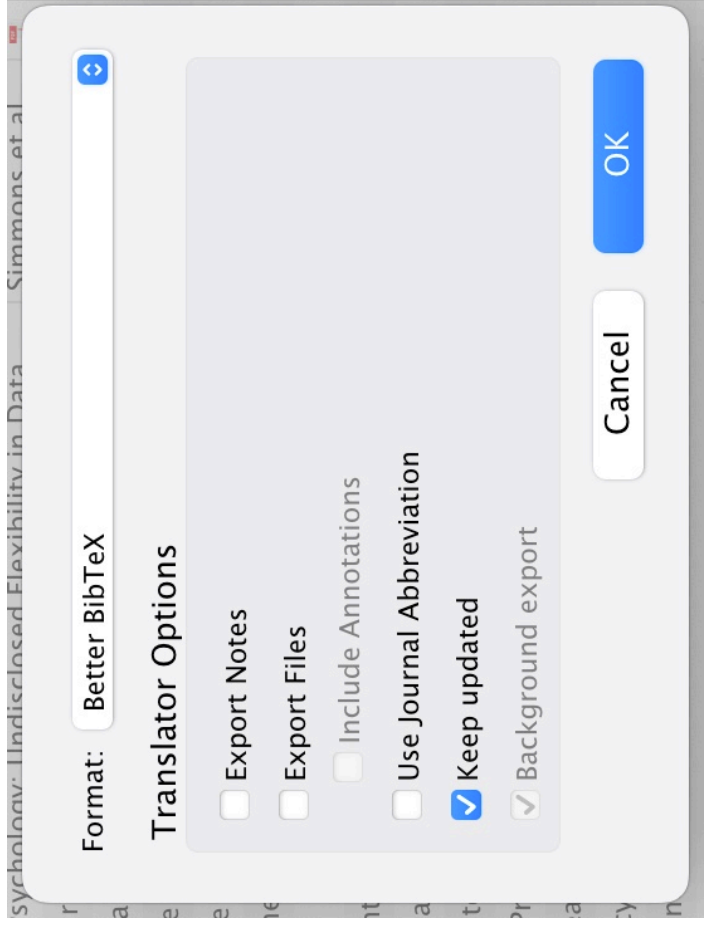
```
---
title: Example report
output:
  pdf_document: paginated
  html_document: chrome_print
bibliography: references.bib
---

# References

::: {refs}
```

Sync Zotero with an RStudio Project

1. Add an Rmd document to your RStudio project called something like `Report.Rmd`
2. Open Zotero and go to `File > Export Library`
 - o Set the format as `Better BibTeX`
 - o Select "Keep updated"
3. Save the file into your RStudio project and call it `references.bib`



Sync Zotero with an RStudio Project

1. Add an Rmd document to your RStudio project called something like `Report.Rmd`
2. Open Zotero and go to `File > Export Library`
 - o Set the format as `Better BibTeX`
 - o Select "Keep updated"
3. Save the file into your RStudio project and call it `references.bib`
4. Open Zotero preferences and go to the `Better BibTeX` section
 - Choose the "Automatic Export" tab
 - Set "Automatic export" to "On Change"



Sync Zotero with an RStudio Project

1. Add an Rmd document to your RStudio project called something like `Report.Rmd`
2. Open Zotero and go to `File > Export library`
3. Save the file into your RStudio project and call it `references.bib`
4. Save the file into your RStudio project and call it `references.bib`
5. Open Zotero preferences and go to the `Better BibTeX` section
6. In your Rmd file cite the paper we added using the citation key... and knit the document!

```

---
title: Example report
output: pdf_document
page: 1
knitpagedown: chrome_print
bibliography: references.bib
---

Ah, yes! The @ioannidis_why_2005

# References

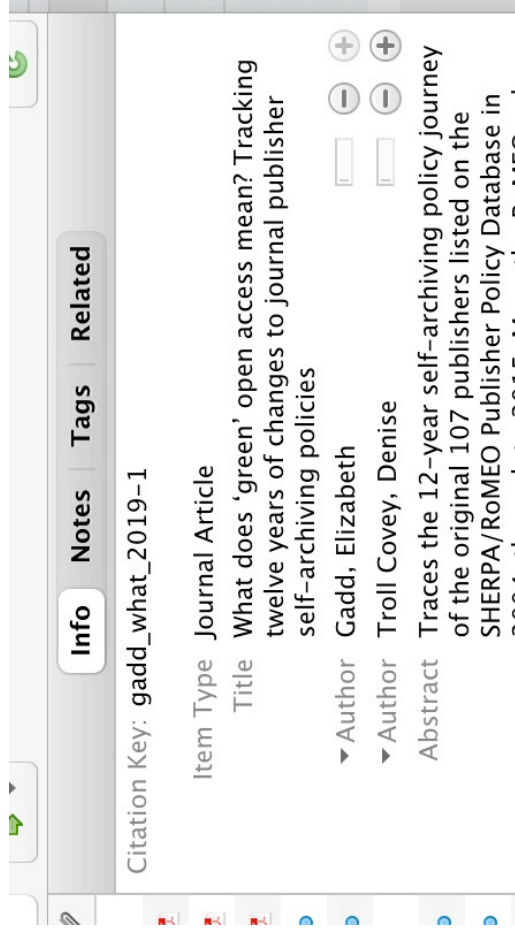
::: # {refs}
:::

```

Zotero citation keys

Zotero automatically generate citation keys for you.

Interestingly again, editing these yourself - they're basically machine-readable data instead of human-readable data.



Challenge

1. Find a paper from a journal
2. Add it to Zotero with the browser plug-in
3. Cite the new paper in your RMarkdown document
4. Knit the .Rmd file together and see the new item in your references

You might want to look at the last few slides of [Week 2's lecture slides](#) which contain at least 40 references.

Zotero and MS Word

Good news Zotero will build automatically install a add-in for citing your collection within MS Word.

1. Open MS Word and look for a Zotero tab. If it's not there continue with these instructions . There are [really good instructions on how to use this add-in](#) on the Zotero web site.
2. Open Zotero preferences and go to Cite > Word Processors
3. Click "Install all Microsoft Word Add-ins"
4. Restart Zotero and MS Word

Adding things that aren't papers to Zotero

It's really easy to use Zotero with papers that have DOI.

But if you want to cite a website or white paper you might want to add this manually to your collection.

- Ensure that all items have a date.
 - Often web pages don't have dates, but check the URL just in case
turning.ac.uk/.../2021-06/...pdf
 - If no date is available go for the current year
- If no authors are listed then use the name of the website or institution
- Add an "access date"

Shared Zotero collection

I've created a shared Zotero collection for this semester -
<https://www.zotero.org/groups/4758381/eng7218-shared-2022/library>.

Let's discuss how we might use this collection as a class.

What reproducibility skills have we learned?

- You need an ORCID! Get one and use it everywhere.
- Use Figshare to release research outputs before you start
 - It'll make it easier to write your research up!
- Use Figshare to store your talks, posters and anything else you want to make available for the future
- Use GitHub for version controlling your code
 - Keep a .gitignore file
- Connect your Figshare and GitHub repos to make your code citable and available for the future
- Use Zotero for storing and managing all your references
- Use the Better BibTeX add-in to keep your bibliography up to date in an RStudio project.

Advanced topics (for the end)

To completely ensure the reproducibility of R code projects you would need to ensure that future users have access to the exact same packages as you used.

The best solution for this is to use the `{renv}` package.

I've personally avoided it for projects recently because I haven't worked on code that I need others to replicate into the future. If you're writing research software then please do consider using it. There's [good documentation](#) for it.